

Foundations of linear regression

Assumptions and construction in Python

- Video:** Make linear regression assumptions 4 min
- Reading:** The four main assumptions of simple linear regression 20 min
- Reading:** Follow-along instructions: Explore linear regression with Python 10 min
- Lab:** Annotated follow-along guide: Explore linear regression with Python 20 min
- Video:** Explore linear regression with Python 9 min
- Reading:** Code functions and documentation 20 min
- Lab:** Activity: Evaluate simple linear regression 1h
- Lab:** Exemplar: Evaluate simple linear regression 20 min
- Practice Quiz:** Test your knowledge: Assumptions and construction in Python 4 questions

Evaluate a linear regression model

Interpret linear regression results

Review: Simple linear regression

Code functions and documentation

In this reading, you will review some of the code from the videos using a different subset of the penguin data. This reading will also share some tips when approaching the statsmodels documentation. This is a good opportunity to review Python functionality in conjunction with exploratory data analysis, basic data cleaning, and model construction.

Review functions from video

Load the dataset

The first few lines of code set up the coding environment and loaded the data. As you might be familiar with, you can call on the `import` function to import any necessary packages. You should use conventional aliases as needed. The example below references a dataset on penguins available through the `seaborn` package.

```
1 # Import packages
2 import pandas as pd
3 import seaborn as sns
4
5 # Load dataset
6 penguins = sns.load_dataset("penguins")
7
8 # Examine first 5 rows of dataset
9 penguins.head()
```

Clean data

After loading the data, the data was cleaned up to create a subset of data for the purposes of our course. The example isolates just the Chinstrap penguins from the dataset and drops rows with missing data.

The index of the dataframe is reset using the `reset_index()` function. When you subset a dataframe, the original row indices are retained. For example, let's say there were Adelie or Gentoo penguins in rows 2 and 3. By subsetting the data just for Chinstrap penguins, your new dataframe would be listed as row 1 and then row 4, as rows 2 and 3 were removed. By resetting the index of the dataframe, the row numbers become rows 1, 2, 3, etc. The data frame becomes easier to work with in the future.

Review the code below. You are encouraged to run the code in your own notebook.

```
1 # Subset just Chinstrap penguins from data set
2 chinstrap_penguins = penguins[penguins["species"] == "Chinstrap"]
3
4 # Reset index of dataframe
5 chinstrap_penguins.reset_index(inplace = True, drop = True)
```

Setup for model construction

Now that the data is clean, you are able to plot the data and construct a linear regression model. First, extract the one X variable, `flipper_length_mm`, and the one Y variable, `bill_depth_mm`, that you are targeting.

```
1 # Subset Data
2 ols_data = chinstrap_penguins[["bill_depth_mm", "flipper_length_mm"]]
```

Because this example is using statsmodels, save the ordinary least squares formula as a string so the computer can understand how to run the regression. The Y variable, `flipper_length_mm` comes first, followed by a tilde and the name for the X variable, `bill_depth_mm`.

```
1 # Write out formula
2 ols_formula = "flipper_length_mm ~ bill_depth_mm"
```

Construct the model

In order to construct the model, you'll first need to import the `ols` function from the `statsmodels.formula.api` interface.

```
1 # Import ols function
2 from statsmodels.formula.api import ols
```

Next, plug in the formula and the saved data into the `ols` function. Then, use the `fit` method to fit the model to the data. Lastly, use the `summary` method to get the results from the regression model.

```
1 # Build OLS, fit model to data
2 OLS = ols(formula = ols_formula, data = ols_data)
3 model = OLS.fit()
4 model.summary()
```

OLS Regression Results					
Dep. Variable:	flipper_length_mm	R-squared:	0.337		
Model:	OLS	Adj. R-squared:	0.327		
Method:	Least Squares	F-statistic:	33.48		
Date:	Wed, 08 Jun 2022	Prob (F-statistic):	2.16e-07		
Time:	13:17:13	Log-Likelihood:	-215.62		
No. Observations:	68	AIC:	435.2		
Df Residuals:	66	BIC:	439.7		
Df Model:	1				
Covariance Type: nonrobust					
	coef	std err	t	P> t	[0.025 0.975]
Intercept	128.6967	11.623	11.073	0.000	105.492 151.902
bill_depth_mm	3.6441	0.630	5.786	0.000	2.387 4.902
Omnibus:	1.350	Durbin-Watson: 1.994			
Prob(Omnibus):	0.509	Jarque-Bera (JB): 0.837			
Skew:	-0.255	Prob(JB): 0.658			
Kurtosis:	3.190	Cond. No. 303.			

Navigating statsmodels documentation

It can require significant work to approach a new Python package or a new set of Python functions, especially when first coding. The benefit of Python being an open source programming language is that there is a strong Python community asking and answering questions. Part of being a successful data professional is knowing how to make your code work and troubleshooting when your code breaks. One way to do this is to go directly to the source, or the official documentation of a particular package.

You've been using the `statsmodels` package to build simple linear regression models. The [statsmodels documentation](#) is available online and is updated regularly. Specifically, you are using the [statsmodels.formula.api interface](#) to perform ordinary least squares estimation.

Examining the page on the `ols` function or the function that performs OLS estimation, you will observe the different function parameters that are allowed, with some notes about each.

Unfortunately, at this time, the statsmodels documentation does not include code examples of how to use the function. If you find documentation that doesn't provide as many examples as you need, or documentation that doesn't provide examples that you need to troubleshoot your code, remember that you can always search online for the function you're trying to use and explore how others in the Python community have handled comparable problems.

Key takeaways

